

The Silent Majority

How Identity Inference Silences Mainstream Opinions

Julian Tait

September 2026

The shy majority



The American context

78% of U.S. adults support building more solar farms.

Source: Pew Research Center, June 2024.

The American context

78% of U.S. adults support building more solar farms.

However, supporters may be **discouraged from voicing it** (Dixon et al. 2024).

Pluralistic ignorance: quiet, because they expect others to differ.

Source: Pew Research Center, June 2024.

The American context

78% of U.S. adults support building more solar farms.

However, supporters may be **discouraged from voicing it** (Dixon et al. 2024).

Pluralistic ignorance: quiet, because they expect others to differ.

Supporters know they are the majority (Dixon et al. 2024), **yet the majority still self-censors** (Burnett et al. 2022).

Source: Pew Research Center, June 2024.

Same opinion with an identity story

**What if it is not the mean norm,
but what it signals about identity?**

78% of U.S. adults support building more solar farms.


Source: Pew Research Center, June 2024.

Same opinion with an identity story

What if it is not the mean norm,
but what it signals about identity?


Republicans

$$\pi_{R, \text{yes}} = 0.64$$



Democrats


$$\pi_{D, \text{yes}} = 0.91$$



Source: Pew Research Center, June 2024.

Same opinion with an identity story

What if it is not the mean norm,
but what it signals about identity?




“Yes” is an **unrepresentative** opinion for a Republican.

Source: Pew Research Center, June 2024. See Bordalo et al. (2016).

Same opinion with an identity story

What if it is not the mean norm,
but what it signals about identity?



“Yes” is an **unrepresentative** opinion for a Republican. $R = 0.64/0.91 \approx 0.70$


Source: Pew Research Center, June 2024. See Bordalo et al. (2016).

The Research Question

How can identity inference contribute to silence?


The Research Question

How can identity inference contribute to silence?



The Research Question

How can identity inference contribute to silence?



Why it matters

When stating an opinion may be costly, we know people:

- ▶ **lie or distort** (Braghieri 2024), **signal another identity** (Lane 2024), **provide a rationale** (Bursztyn et al. 2023), . . .

Why it matters

When stating an opinion may be costly, we know people:

- ▶ **lie or distort** (Braghieri 2024), **signal another identity** (Lane 2024), **provide a rationale** (Bursztyn et al. 2023), . . . or **say nothing**.

Why it matters

When stating an opinion may be costly, we know people:

- ▶ **lie** or distort (Braghieri 2024), **signal another identity** (Lane 2024), **provide a rationale** (Bursztyn et al. 2023), . . . or **say nothing**.
- ▶ silence **leaves no trace** (Noelle-Neumann 1974; Ho & Huang 2026).

Why it matters

When stating an opinion may be costly, we know people:

- ▶ lie or distort (Braghieri 2024), **signal another identity** (Lane 2024), **provide a rationale** (Bursztyn et al. 2023), . . . or **say nothing**.
- ▶ silence **leaves no trace** (Noelle-Neumann 1974; Ho & Huang 2026).

The cost is social: false polarisation (Bordalo et al. 2021) · norms the private majority rejects persist (Bursztyn et al. 2020) · stalls collective action (Andre et al. 2024).

Why it matters


When stating an opinion may be costly, we know people:

- ▶ lie or distort (Braghieri 2024), **signal another identity** (Lane 2024), **provide a rationale** (Bursztyn et al. 2023), . . . or **say nothing**.
- ▶ silence **leaves no trace** (Noelle-Neumann 1974; Ho & Huang 2026).

The cost is social: **false polarisation** (Bordalo et al. 2021) · **norms the private majority rejects persist** (Bursztyn et al. 2020) · **stalls collective action** (Andre et al. 2024).

Can identity inference alone silence even a majority view?

Abstracting



Real opinions carry **priors**, **moral** weight, **entrenched identities** and **second-order beliefs**.

Abstracting


Triangles

$$\pi_{\Delta, \text{green}} = 0.64$$



Circles

$$\pi_{\circ, \text{green}} = 0.91$$



$$\text{identity} \in \{ \Delta, \circ \}$$

$$\text{opinion} \in \{ \text{green blob}, \text{blue blob} \}$$

The game

The aim: see what people say when their words may classify them.

The game

The aim: see what people say when their words may classify them.


Stage 0

senders receive an
identity and an opinion




The game

The aim: see what people say when their words may classify them.




The game

The aim: see what people say when their words may classify them.



The game

The aim: see what people say when their words may classify them.




- ▶ Senders are **paid to share** their opinion: +2 to most (70%), -2 to the rest (30%).
- ▶ Observers are **paid to guess** the sender's identity correctly (+ conf.).
- ▶ **Treatment.** STAKES: being misclassified costs you. NOSTAKES: no cost.

Design

Vary the **diagnosticity** of opinion 1 (how strongly it points to one identity), while:


- ▶ fixing the **identity split**, at $\Pr(A) = \frac{1}{2}$;
- ▶ fixing the **cost** of misclassification.

SYMMETRIC



$$P(A | 1) = 0.50$$

MAJORITY




$$P(A | 1) \approx 0.41, R(1, A) \approx 0.68$$

Design

H1. Silence rises when opinion 1 points to the **other identity**: **MAJORITY** > SYMMETRIC.


H2. The gap is **inference**: present under STAKES, gone under NOSTAKES.

SYMMETRIC



$$P(A | 1) = 0.50$$

MAJORITY



$$P(A | 1) \approx 0.41, R(1, A) \approx 0.68$$

Prediction: classification & confidence

Exploratory.

- ▶ Observer posteriors deviate from Bayes toward representativeness: confidence over-shoots $P(\text{identity} \mid \text{opinion})$ (à la Bordalo).

Summary

- ▶ The within-identity **majority** can be the suppressed opinion.
- ▶ Suppression is about the **inferred identity**, not just the mean norm.
- ▶ Suppression **survives full information**: subjects are given full information about the distribution.

Thank you.

Discussion

- ▶ **Pricing expression directly?**

- ▶ How much we pay them to share the opinion: the price of expression.
- ▶ How much they pay to send the message (to identify): how much their identity matters to them.
- ▶ Here: binary share-or-silence, so a silence rate, not a price.

- ▶ **Alternative story: why majorities DO protect.**

- ▶ People feel protected by lower encounter rates with “others”.
- ▶ Study 2 (real-world identities/opinions) can build it in through in-group favouritism of opinions.

Real-world version


	Republicans	Democrats
First-trimester abortion legal	59%	92%
Climate change is happening	~60%	~90%



Source: Gallup 2024; Pew Research Center, 2023.

Real-world version

	Republicans	Democrats
First-trimester abortion legal	59%	92%
Climate change is happening	~60%	~90%




Both are elicited, not set by us.

Source: Gallup 2024; Pew Research Center, 2023.


The MINORITY comparison

The design includes a third treatment. **MINORITY** and **MAJORITY** hold diagnosticity $R(1, A)$ fixed but sit on opposite sides of the identity split: opinion 1 is the within-identity **majority** view in one and the **minority** view in the other.

MAJORITY



MINORITY



$$R(1, A) \approx 0.68, P(A | 1) \approx 0.41$$

$$R(1, A) \approx 0.69, P(A | 1) \approx 0.41$$

$R(1, A)$ and the naive Bayes posterior are matched across the pair.

The MINORITY comparison: H3


Across the two matched treatments, opinion 1 is the within-identity **majority** view in **MAJORITY** and the **minority** view in **MINORITY**, with diagnosticity $R(1, A)$ and the naive Bayes posterior held **equal**.

H3. A within-identity majority does not switch off identity inference:

$$\text{MAJORITY} \approx \text{MINORITY} \quad (\text{matched } R(1, A)).$$

Because R is matched across the pair, any difference in silence between them cannot come from how diagnostic the opinion is. Equality is the prediction if suppression runs off the *inferred identity* rather than the raw split.

The MINORITY comparison: robustness



The gap stays positive across the design's own regime; it closes only outside it (high baseline silence, or group A a population minority).

Dixon 2024 vs. Burnett 2022

	Dixon et al. 2024	Burnett et al. 2022
Who	1000 Republicans only ; YouGov, Dec 2022. Majority/minority is <i>within</i> Republicans on climate policy: supporters = majority, opposers = minority	1047 US social-media users; Qualtrics, even R/D split
Measures	willingness to speak about climate to other Republicans, 1–6 ($M = 2.25$, lower = more willing); no speech or silence observed	willingness-to-self-censor scale; “majority” is compositional ($\approx 64\%$ self-placing in the ideological middle band, grouped by the authors); no issue opinion measured
Finding	supporters perceive their majority correctly; a perceived minority-dominated environment is associated, via anticipated conflict, with marginally lower willingness (tiny but significant)	the moderate middle self-censors at least as much as strong liberals (overall effect significant, small)
Cannot show	that knowers fall silent : the effect is marginal and the level still leans willing	knowing, a majority of <i>beliefs</i> on any question, or observed silence

Marginal vs level: Dixon supports discouraged-*at-the-margin*, not fall-silent — the level still leans willing ($M = 2.25/6$). **Neither** paper shows knowing *and* silence in the same people; our design induces exactly that.

Bursztyn et al. 2020 vs. our model

		Bursztyn et al. 2020 (Saudi WWOH)	Ours
Stage 1 (sender)		binary conviction $t \in \{A, B\}$ (pro/anti women working); picks observable action $s \in \{A, B\}$, may misreport at lying cost c_i	identity g and opinion t ; shares the opinion or stays silent (no lying)
Stage 2 (receiver)		posterior $\Pr(t = A \mid s)$; image payoff affine in it	guesses modal group $\hat{g}$ from the message
Type / inference		1D: conviction only; scalar posterior	2D : group $\times$ opinion; classification via cross-group ratio $R(t, g)$
Treatment		shifts belief q about the audience; joint fixed	shifts the joint π structurally (hold $\pi_{1,A}$, move $\pi_{1,B}$)
Correcting beliefs		moves behaviour	within-group share does not ; only cross-group R matters

Overlap: both are sender–receiver signalling games with a binary type and a reputational image term, and both correct a misperception.

Ours adds: the 2D structure separates within-group share from cross-group diagnosticity, so a within-group **majority** opinion can still be suppressed (impossible in 1D); and the message is **silence**, not a lie.

Cited work: what each paper measures (1/2)

Braghieri 2024. A lab-in-field experiment (US students) reporting sensitive attitudes in private vs public. Public agreement shifts ~ 0.13 SD toward the acceptable pole and public statements carry less information — self-reported attitudes in a reporting task, not observed speech.

Lane 2024. A lab experiment where subjects choose which identity dimension is salient before economic play. Many sacrifice payoff to foreground a preferred identity; it is about *which identity to reveal*, not suppressing a majority opinion.

Bursztyn et al. 2023. Survey experiments randomising whether a credible rationale gives social cover before a stigmatised post. Cover raises willingness to dissent (+11–17 pp) and lowers inferred prejudice; the outcome is an incentivised posting decision with real stakes, not naturally observed speech.

Noelle-Neumann 1974. The spiral of silence: fearing isolation, those who feel in the minority withdraw, making the view look weaker still. A foundational theory paper with illustrative Allensbach survey data, driven by *perceived minority status* — distinct from identity inference.

Ho & Huang 2026. A model of self-censorship plus limited attention to silence, tested in a dynamic lab discussion. Self-censorship of inappropriate views amplifies over time and raising attention to silence improves belief accuracy — the driver is *perceived social (in)appropriateness*.

Cited work: what each paper measures (2/2)

Bordalo et al. 2016. A representativeness-heuristic model of group beliefs: distinctively-common types (high likelihood ratio R) are overweighted, exaggerating group differences. It models *belief formation* only and supplies the R machinery we borrow — no expression or silence decision.

Bordalo et al. 2021. An observational ANES study ($\sim 10,500$ partisans, 1980–2004) of perceived vs actual party positions. Perceived partisan differences exaggerate actual ones by $\sim 70\%$, more on salient issues — self-reported beliefs, not behaviour.

Bursztyn et al. 2020. A field experiment with Saudi men: correct their underestimate of other men's support for women working, then measure a real job-search sign-up. Correcting beliefs *about others* moved behaviour — a one-dimensional norm game, opposite in direction to our within-group claim.

Andre et al. 2024. A globally representative survey, 125 countries, $\sim 130,000$ respondents. A large majority support climate action, yet people underestimate others' willingness by ~ 26 pp (69% vs 43%) — the canonical pluralistic-ignorance statistic (self-reported perceptions).

Dixon et al. 2024. A cross-sectional YouGov survey, 1,000 US Republicans, with a 6-point willingness-to-speak scale. Supporters correctly know they are the majority, but willingness leans *willing* ($M = 2.25/6$) with no silent share — correlational.

Burnett et al. 2022. A cross-sectional survey ($N = 1,047$) using a Willingness-to-Self-Censor scale, a moderate “majority” vs ideological extremes. The middle self-censors slightly more ($\eta^2 = .02$); there is *no perception measure*, and the scale is a disposition, not observed silence.