

The Silent Majority

How Identity Inference Silences Mainstream Opinions

Julian Tait

September 2026

The shy majority



The American context

78% of U.S. adults support building more solar farms.

Source: Pew Research Center, June 2024.

The American context

78% of U.S. adults support building more solar farms.

However, supporters may be **discouraged from voicing it** (Dixon et al. 2024).

Source: Pew Research Center, June 2024.

The American context

78% of U.S. adults support building more solar farms.

However, supporters may be **discouraged from voicing it** (Dixon et al. 2024).

Pluralistic ignorance: quiet, because they expect others to differ.

Source: Pew Research Center, June 2024.

The American context

78% of U.S. adults support building more solar farms.

However, supporters may be **discouraged from voicing it** (Dixon et al. 2024).

Pluralistic ignorance: quiet, because they expect others to differ.

Large pluralistic ignorance in climate (Andre et al. 2024).

Source: Pew Research Center, June 2024.

The American context

78% of U.S. adults support building more solar farms.

However, supporters may be **discouraged from voicing it** (Dixon et al. 2024).

Pluralistic ignorance: quiet, because they expect others to differ.

Large pluralistic ignorance in climate (Andre et al. 2024).

Supporters know they are the majority (Dixon et al. 2024), **yet the majority still self-censors** (Burnett et al. 2022).

Source: Pew Research Center, June 2024.

Same opinion with an identity story

**What if it is not the mean norm,
but what it signals about identity?**

78% of U.S. adults support building more solar farms.


Source: Pew Research Center, June 2024.

Same opinion with an identity story

What if it is not the mean norm,
but what it signals about identity?


Republicans

$$\pi_{R, \text{yes}} = 0.64$$



Democrats


$$\pi_{D, \text{yes}} = 0.91$$



Source: Pew Research Center, June 2024.

Same opinion with an identity story

What if it is not the mean norm,
but what it signals about identity?



“Yes” is an **unrepresentative** opinion for a Republican.


Source: Pew Research Center, June 2024. See Bordalo et al. (2016).

The Research Question

How can identity inference contribute to silence?


The Research Question

How can identity inference contribute to silence?



The Research Question

How can identity inference contribute to silence?



Why it matters I

When an opinion is costly to state, people have options.

- ▶ **Lie**, or distort (Braghieri 2024).
- ▶ **Signal another identity** (Lane 2024).
- ▶ **Give a rationale** (Bursztyn et al. 2023).
- ▶ **Say nothing.**

Why it matters I

When an opinion is costly to state, people have options.

- ▶ **Lie**, or distort (Braghieri 2024).
- ▶ **Signal another identity** (Lane 2024).
- ▶ **Give a rationale** (Bursztyn et al. 2023).
- ▶ **Say nothing.**

Only silence **leaves no trace** (Noelle-Neumann 1974; Ho & Huang 2026).

Why it matters I

When an opinion is costly to state, people have options.

- ▶ **Lie**, or distort (Braghieri 2024).
- ▶ **Signal another identity** (Lane 2024).
- ▶ **Give a rationale** (Bursztyn et al. 2023).
- ▶ **Say nothing.**

Only silence **leaves no trace** (Noelle-Neumann 1974; Ho & Huang 2026).

Opinion is **not missing at random**: it is selected on identity.

Why it matters II

Belief correction cannot fix it: their beliefs are **already correct**.

Why it matters II

Belief correction cannot fix it: their beliefs are **already correct**.

The cost of the silence. It is **widespread**: **62%** of Americans hold political views they are afraid to share (Ekins/Cato 2020). At scale it:

- drives **false polarisation** — perceived partisan gaps run $\sim 70\%$ larger than the real ones (Bordalo et al. 2021);
- lets a norm the **private majority rejects** persist — 82% privately approve, yet stay quiet (Bursztyn et al. 2020);
- stalls **collective action** as majorities underestimate their own support (Andre et al. 2024, 125 countries).

Why it matters II


Belief correction cannot fix it: their beliefs are **already correct**.

The cost of the silence. It is **widespread**: **62%** of Americans hold political views they are afraid to share (Ekins/Cato 2020). At scale it:

- drives **false polarisation** — perceived partisan gaps run $\sim 70\%$ larger than the real ones (Bordalo et al. 2021);
- lets a norm the **private majority rejects** persist — 82% privately approve, yet stay quiet (Bursztyn et al. 2020);
- stalls **collective action** as majorities underestimate their own support (Andre et al. 2024, 125 countries).

Ours: how **identity inference** silences a within-identity **majority opinion** — suppression that runs off the **inferred identity**, not the mean norm, and survives full information.

Abstracting



Real opinions carry **priors**, **moral** weight and **entrenched identities**.

Abstracting


Triangles

$$\pi_{\Delta, \text{green}} = 0.64$$



Circles

$$\pi_{\circ, \text{green}} = 0.91$$



$$\text{identity} \in \{ \Delta, \circ \}$$

$$\text{opinion} \in \{ \text{green blob}, \text{blue blob} \}$$

The game

Aim: test how strategically people act when their perception matters.


Stage 0

senders receive an
identity and an opinion




The game

Aim: test how strategically people act when their perception matters.



The game


Aim: test how strategically people act when their perception matters.



- ▶ **Paid to share:** sharing is worth +2 to most senders (70%), -2 to the rest (30%).
- ▶ **Treatment.** STAKES: being misclassified costs you. NOSTAKES: no cost.

The game

Aim: test how strategically people act when their perception matters.




- ▶ **Paid to share:** sharing is worth +2 to most senders (70%), -2 to the rest (30%).
- ▶ **Treatment.** STAKES: being misclassified costs you. NOSTAKES: no cost.
- ▶ Observers are paid to guess the sender's identity correctly.

Design

Vary the **diagnosticity** of opinion 1 (how strongly it points to one identity), while:

- ▶ fixing the **identity split**, at $\Pr(A) = \frac{1}{2}$;
- ▶ fixing the **cost** of misclassification.




$$P(A \mid 1) = 0.50$$

Design

Vary the **diagnosticity** of opinion 1 (how strongly it points to one identity), while:


- ▶ fixing the **identity split**, at $\Pr(A) = \frac{1}{2}$;
- ▶ fixing the **cost** of misclassification.

SYMMETRIC



$$P(A | 1) = 0.50$$


MAJORITY



$$R(1, A) \approx 0.68, P(A | 1) \approx 0.41$$


Prediction: staying silent

- **H1.** Silence rises when opinion 1 is held by the **wrong identity**:
MAJORITY > **SYMMETRIC**.



Prediction: staying silent

- ▶ **H1.** Silence rises when opinion 1 is held by the **wrong identity**:
MAJORITY > **SYMMETRIC**.
- ▶ **H2.** The gap is **inference**:
present under **STAKES**, gone under **NOSTAKES**.



Prediction: classification & confidence

Exploratory.

- ▶ Observer posteriors deviate from Bayes toward representativeness: confidence over-shoots $P(\text{identity} \mid \text{opinion})$ (à la Bordalo).

Summary

- ▶ The within-identity **majority** can be the suppressed opinion.
- ▶ Suppression is about the **inferred identity**, not just the mean norm.
- ▶ Suppression **survives full information**: subjects are given full information about the distribution.

Thank you.

Discussion

- ▶ **Pricing expression directly?**
 - ▶ How much we pay them to share the opinion: the price of expression.
 - ▶ How much they pay to send the message (to identify): how much their identity matters to them.
 - ▶ Here: binary share-or-silence, so a silence rate, not a price.
- ▶ **Alternative story: why majorities DO protect.**
 - ▶ People feel protected by lower encounter rates with “others”.
 - ▶ Study 2 (real-world identities/opinions) can build it in through in-group favouritism of opinions.

Real-world version


	Republicans	Democrats
First-trimester abortion legal	59%	92%
Climate change is happening	~60%	~90%



Source: Gallup 2024; Pew Research Center, 2023.

Real-world version

	Republicans	Democrats
First-trimester abortion legal	59%	92%
Climate change is happening	~60%	~90%




Both are elicited, not set by us.

Source: Gallup 2024; Pew Research Center, 2023.


The MINORITY comparison

The design includes a third treatment. **MINORITY** and **MAJORITY** hold diagnosticity $R(1, A)$ fixed but sit on opposite sides of the identity split: opinion 1 is the within-identity **majority** view in one and the **minority** view in the other.

MAJORITY



MINORITY



$$R(1, A) \approx 0.68, P(A | 1) \approx 0.41$$

$$R(1, A) \approx 0.69, P(A | 1) \approx 0.41$$

$R(1, A)$ and the naive Bayes posterior are matched across the pair.

The MINORITY comparison: H3


Across the two matched treatments, opinion 1 is the within-identity **majority** view in **MAJORITY** and the **minority** view in **MINORITY**, with diagnosticity $R(1, A)$ and the naive Bayes posterior held **equal**.

H3. A within-identity majority does not switch off identity inference:

$$\text{MAJORITY} \approx \text{MINORITY} \quad (\text{matched } R(1, A)).$$

Because R is matched across the pair, any difference in silence between them cannot come from how diagnostic the opinion is. Equality is the prediction if suppression runs off the *inferred identity* rather than the raw split.

The MINORITY comparison: robustness



The gap stays positive across the design's own regime; it closes only outside it (high baseline silence, or group A a population minority).

Dixon 2024 vs. Burnett 2022

	Dixon et al. 2024	Burnett et al. 2022
Who	1000 Republicans only ; YouGov, Dec 2022. Majority/minority is <i>within</i> Republicans on climate policy: supporters = majority, opposers = minority	1047 US social-media users; Qualtrics, even R/D split
Measures	willingness to speak about climate to other Republicans, 1–6 ($M = 2.25$, lower = more willing); no speech or silence observed	willingness-to-self-censor scale; “majority” is compositional ($\approx 64\%$ self-placing in the ideological middle band, grouped by the authors); no issue opinion measured
Finding	supporters perceive their majority correctly; a perceived minority-dominated environment is associated, via anticipated conflict, with marginally lower willingness (tiny but significant)	the moderate middle self-censors at least as much as strong liberals (overall effect significant, small)
Cannot show	that knowers fall silent : the effect is marginal and the level still leans willing	knowing, a majority of <i>beliefs</i> on any question, or observed silence

Marginal vs level: Dixon supports discouraged-*at-the-margin*, not fall-silent — the level still leans willing ($M = 2.25/6$). **Neither** paper shows knowing *and* silence in the same people; our design induces exactly that.

Bursztyn et al. 2020 vs. our model

	Bursztyn et al. 2020 (Saudi WWOH)	Ours
Stage 1 (sender)	binary conviction $t \in \{A, B\}$ (pro/anti women working); picks observable action $s \in \{A, B\}$, may misreport at lying cost c_i	identity g and opinion t ; shares the opinion or stays silent (no lying)
Stage 2 (receiver)	posterior $\Pr(t = A \mid s)$; image payoff affine in it	guesses modal group $\hat{g}$ from the message
Type / inference	1D: conviction only; scalar posterior	2D : group $\times$ opinion; classification via cross-group ratio $R(t, g)$
Treatment	shifts belief q about the audience; joint fixed	shifts the joint π structurally (hold $\pi_{1,A}$, move $\pi_{1,B}$)
Correcting beliefs	moves behaviour	within-group share does not ; only cross-group R matters

Overlap: both are sender–receiver signalling games with a binary type and a reputational image term, and both correct a misperception.

Ours adds: the 2D structure separates within-group share from cross-group diagnosticity, so a within-group **majority** opinion can still be suppressed (impossible in 1D); and the message is **silence**, not a lie.